

Proposes *witness generation* procedures for proving program transformations [1]. The main insight is that heuristic passes can produce evidence to help the verifier.

From high-level to low-level, the paper's connection to real life is:

- Compiler optimizations are program transformations.
- A program transformation is correct if running the result (target) gives an output that the source code could have produced.
- Formally, target code T *implements* source code S according to a relation α iff α relates all initial target states with initial source states and whenever the target reduces to a final state, the source can reduce to a final state related by α . Note that this definition ignores intermediate steps. An optimizing pass that produced T given S is correct if we can give a relation α with the above property.
- The paper argues that *stuttering simulations* are a useful class of relations because they are straightforward to prove and imply *implements* relations. If S and T are in a stuttering simulation, then T implements S .
- A stuttering simulation R must:
 - Relate initial target states to initial source states.
 - Provide a well-founded measure $<$ such that, for all transitions $t_1 \rightarrow t_2$ in T and all source states s_1 related by R to t_1 , one of the following must hold:
 - * There is a source state s_2 such that $s_1 \rightarrow s_2$ and $(t_2, s_2) \in R$
 - * We stutter in T , which means $(t_2, s_1) \in R$ and $(t_2, s_1) < (t_1, s_1)$ holds.
 - * We stutter in S ; meaning is a s_2 such that $(t_2, s_2) \in R$ and $(t_2, s_2) < (t_1, s_2)$.

The stuttering means we don't move to a new state along both axis, but instead decrease along $<$.
- Stuttering simulations compose, just like passes compose into a whole compiler.
- Stuttering simulation can prove interesting optimizations. The authors show examples of constant propagation, dead code elimination, CFG compression, loop hoisting, and loop reordering.

Strengths

- It's exciting to hope that we could validate all compiler passes in a uniform way. It's also exciting that we don't need to rewrite old optimizing transformations to do so (we only need to modify them to produce evidence).

Weaknesses

- I'm not certain how the wellfounded relations $<$ compose. Do we get to reset the measure between passes?
- The examples are explained well, but I'm looking forward to a more general way of producing these relations during an optimization pass.

References

- [1] Kedar S. Namjoshi and Lenore D. Zuck. Witnessing program transformations. In *SAS*, 2013.