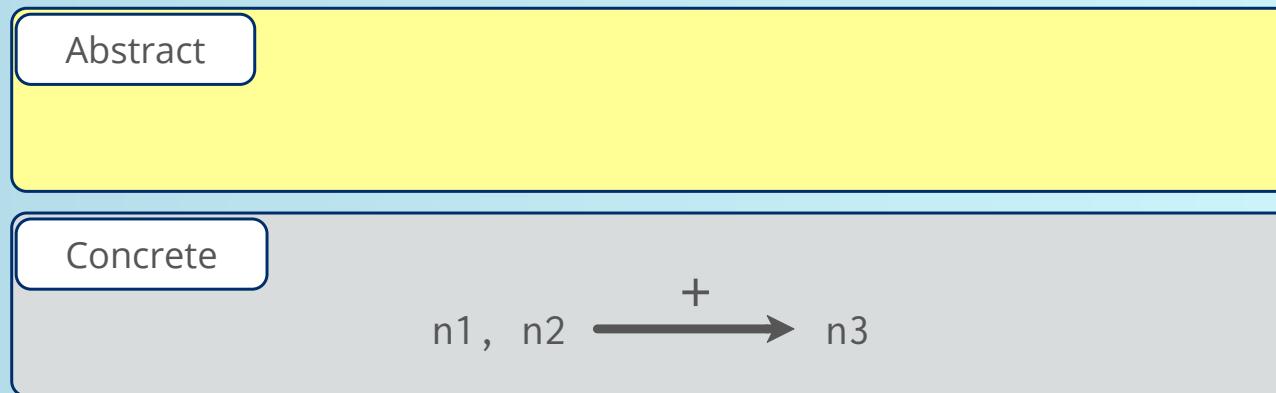


Nice to Meet You: Synthesizing Practical MLIR Abstract Transformers

Xuanyu Peng Dominic Kennedy Yuyou Fan
Ben Greenman John Regehr Loris D'Antoni



Problem: Finding Abstract Transformers $f^\#$



Problem: Finding Abstract Transformers $f^\#$

Abstract: Interval Analysis

$$[l_1, u_1], [l_2, u_2] \xrightarrow{+^\#} [l_3, h_3]$$

Concrete

$$n_1, n_2 \xrightarrow{+} n_3$$

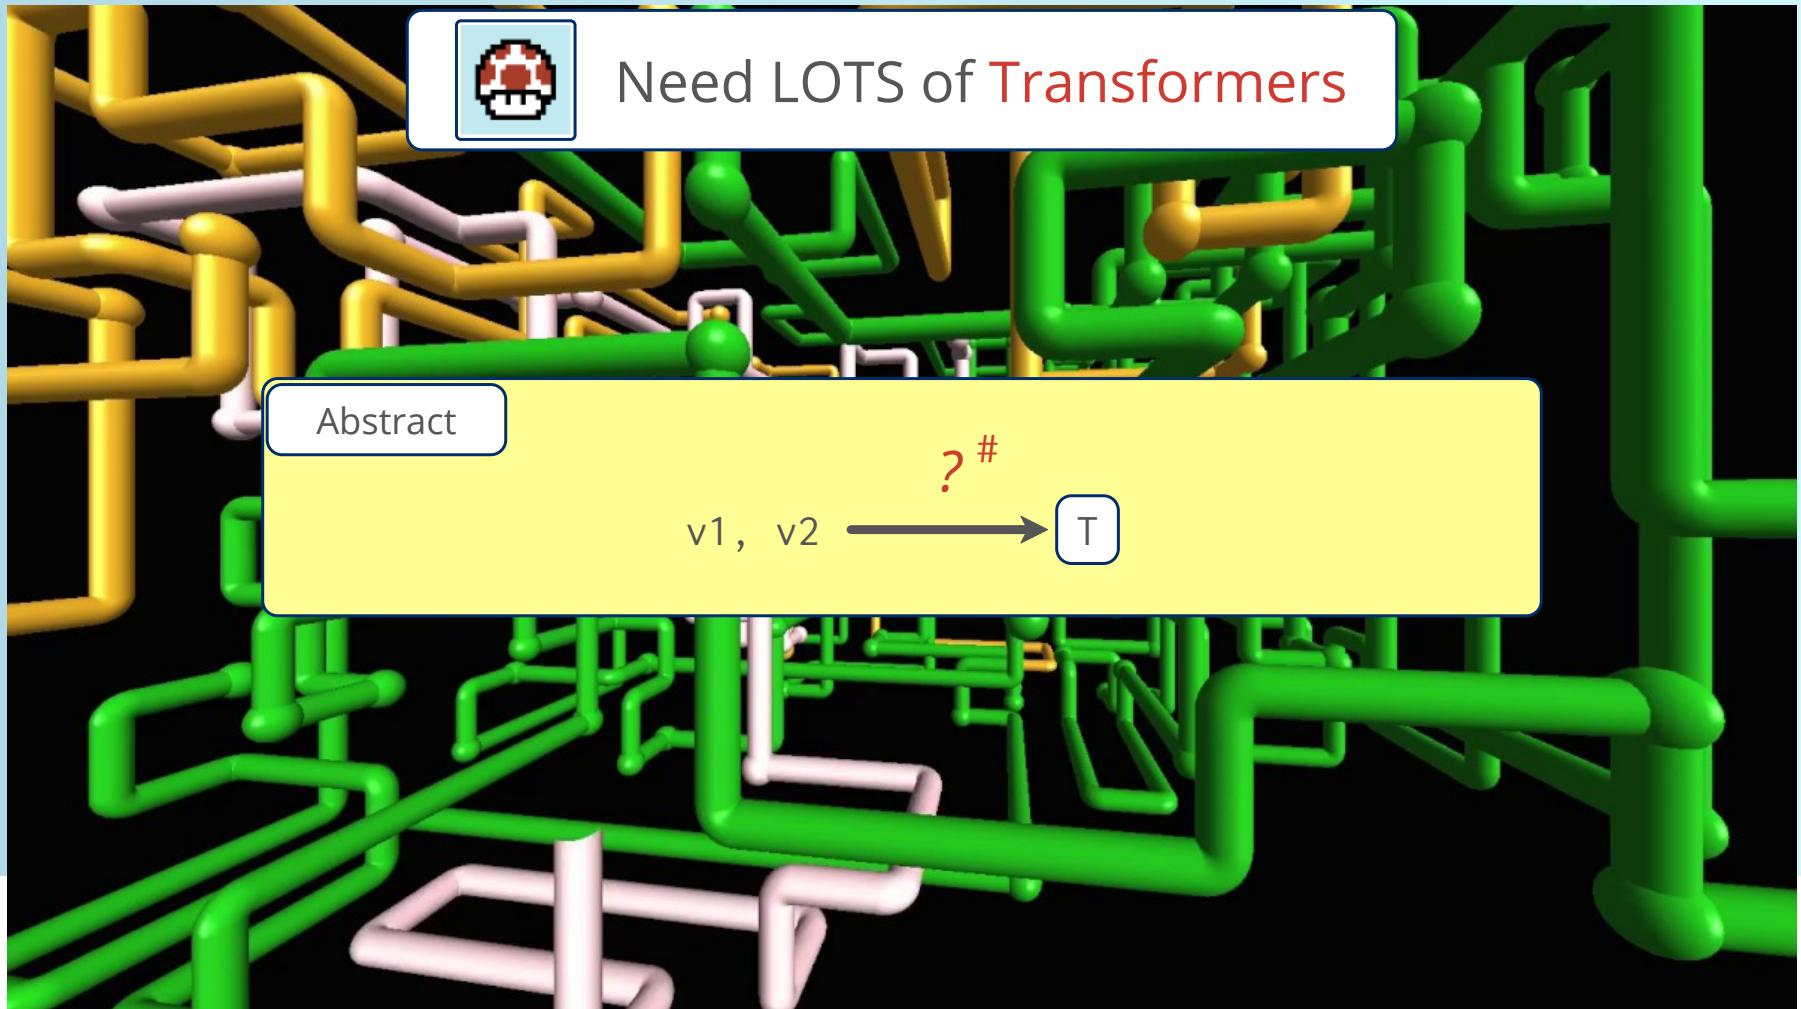
Problem: Finding Abstract Transformers $f^\#$

Abstract: Interval Analysis

$$[11, u1], [12, u2] \xrightarrow{+, \text{div}, \max, \dots} [13, h3]$$







Nick Lewycky 2015-03-24 18:46:31 PDT

[Description](#)

```
$ clang++ -v
clang version 3.7.0 (trunk 233044)
Target: x86_64-unknown-linux-gnu
```

Testcase:

```
#include <stdio.h>
#include <stdlib.h>
#include <string>

using namespace std;

int main(int argc, char **argv) {
    int r = 2;
    bool ok = true;
    while (ok) {
        string ab;
        for (int i = 0; i < r % 3; i++) {
            ab += "ab";
        }
        printf("%d %s\n", r, ab.c_str());
        r++;
        ok = (r < 3);
    }
}
```

```
nlewycky@ducttape:~$ clang++ -O2 a.cc -o a
nlewycky@ducttape:~$ ./a
2 ab
nlewycky@ducttape:~$ clang++ a.cc -o a
nlewycky@ducttape:~$ ./a
2 abab
```

The -O0 result is correct.

[Bug 23011 - miscompile of % in loop](#)

[Status: RESOLVED FIXED](#)

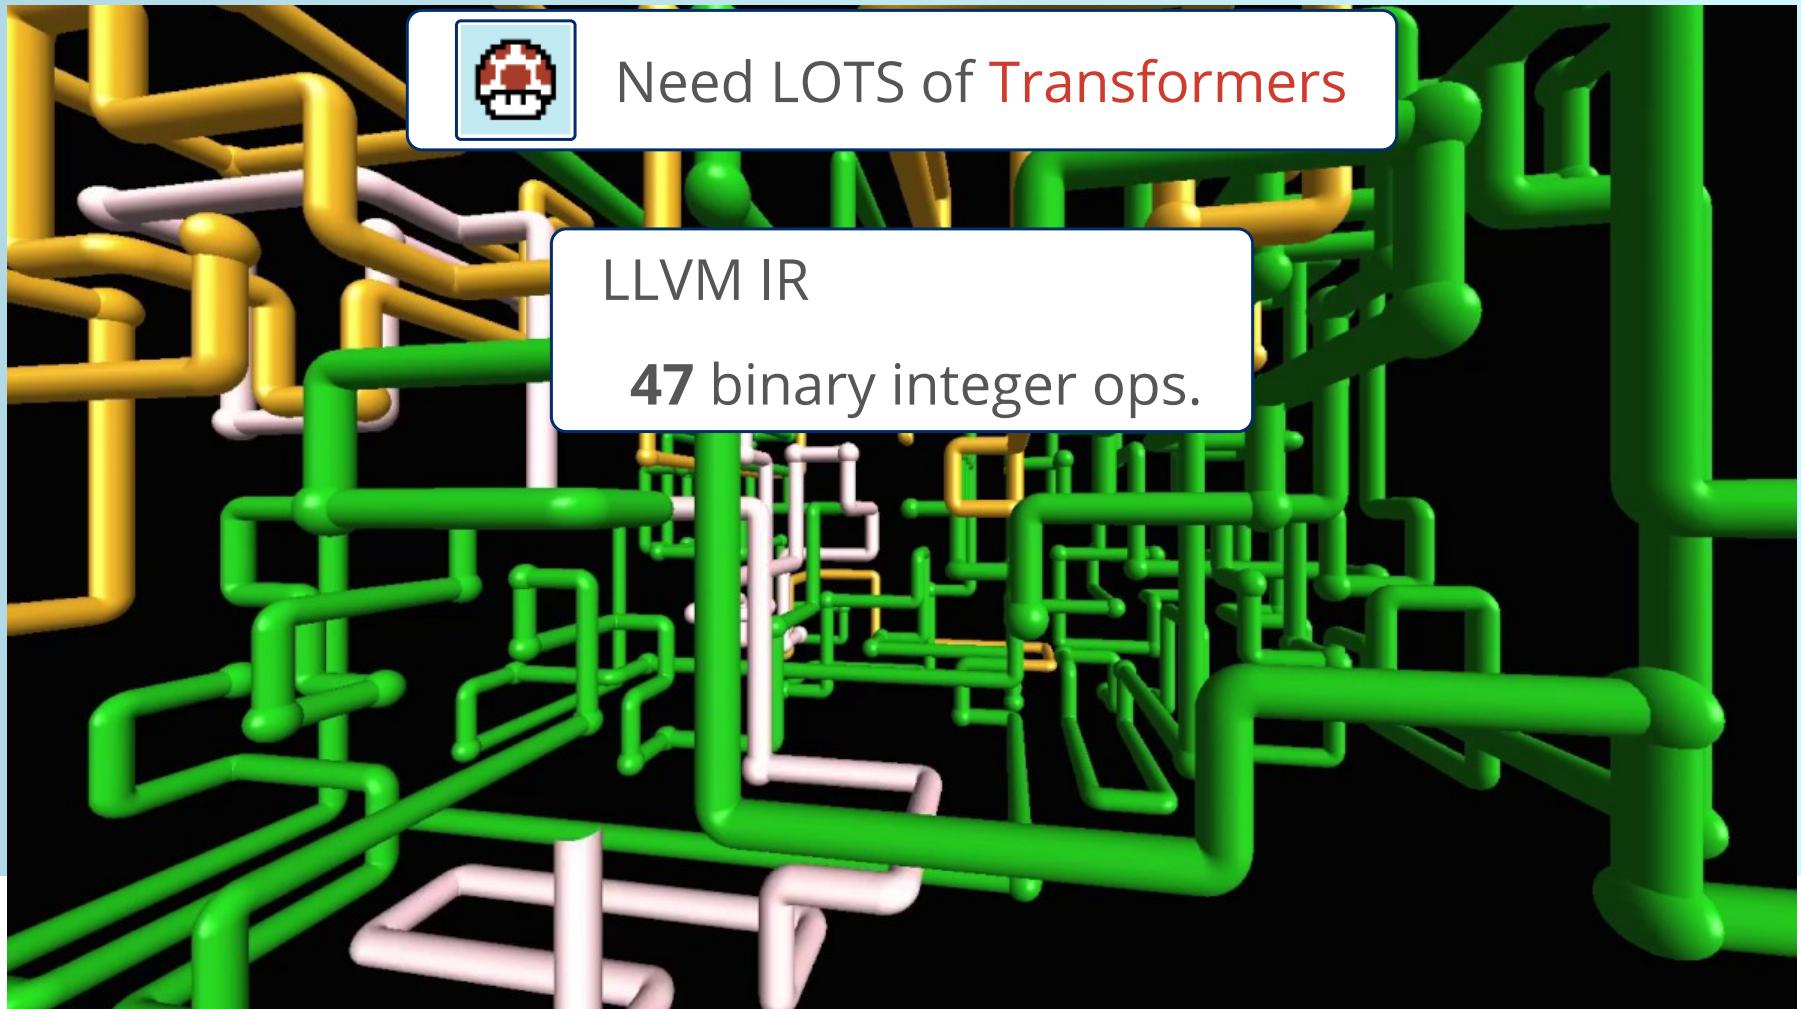
Sanjoy Das 2015-03-24 21:49:17 PDT

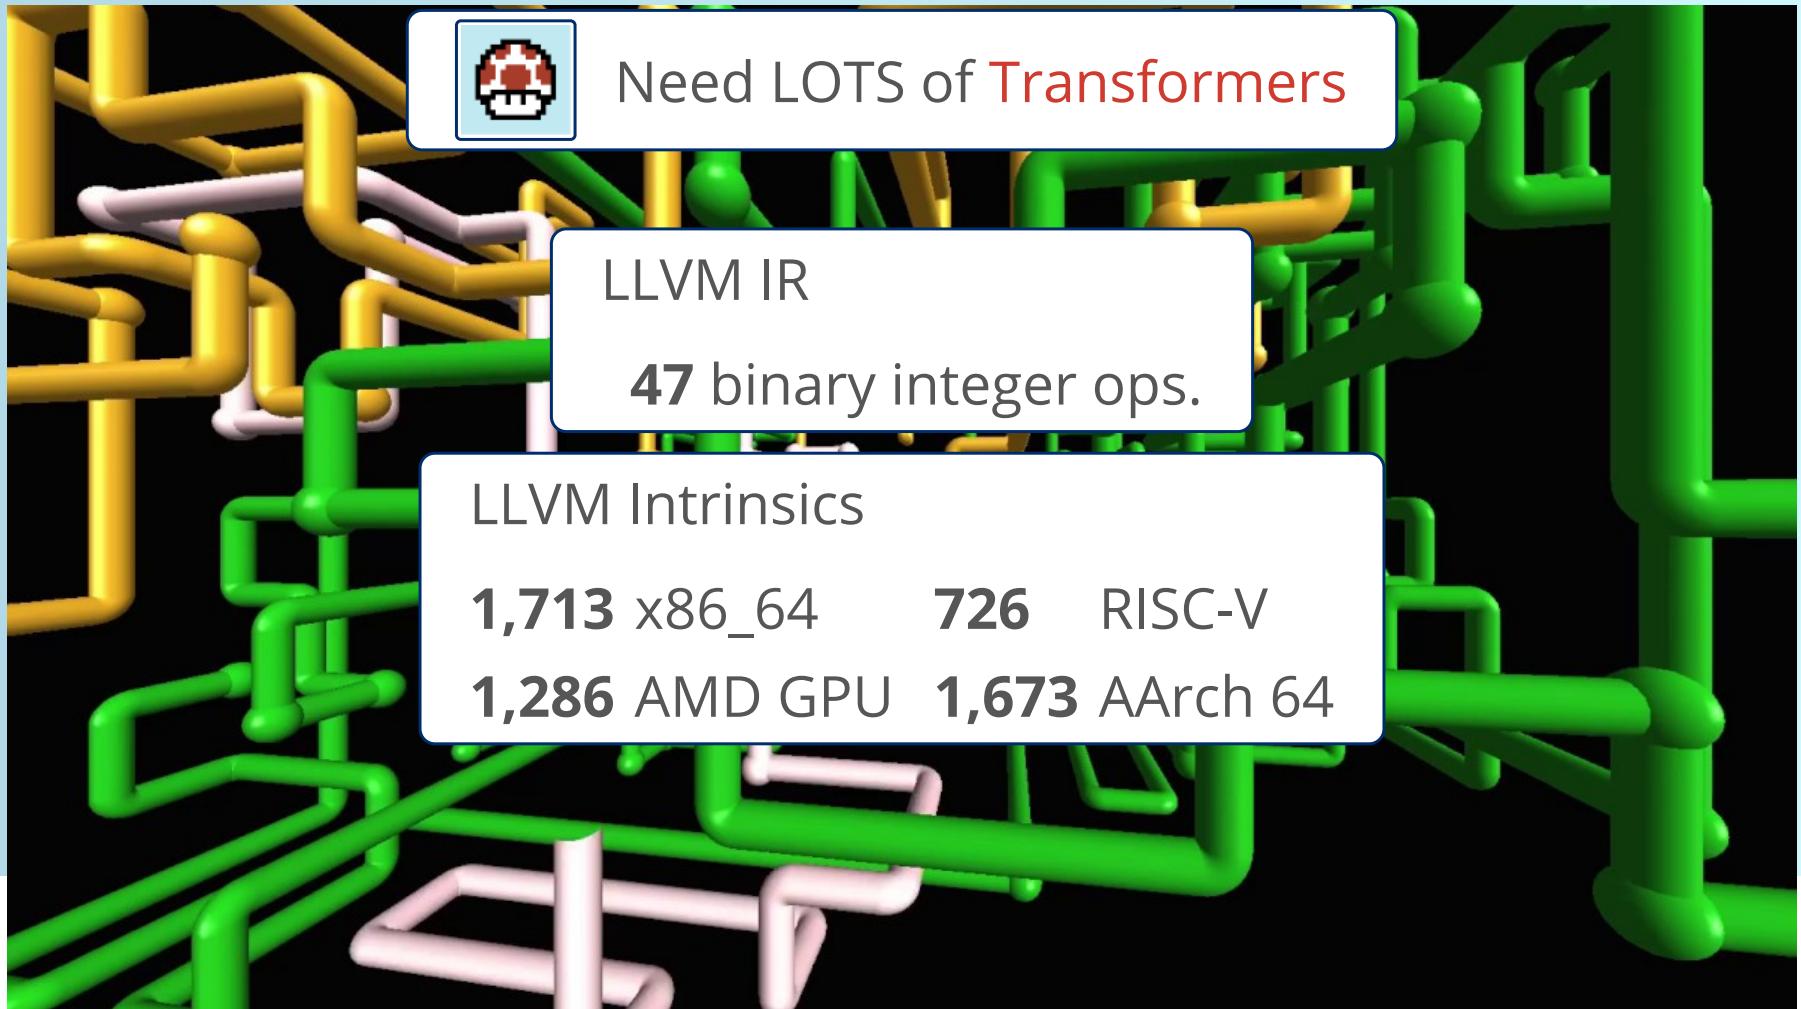
[Comment 5](#)

ComputeNumSignBits(%rem) returns 31, should be 30.
AFAICT, the bug is in ValueTracking:

```
// Calculate the leading sign bit constraints by examining the
// denominator. The remainder is in the range 0..C-1, which is
// calculated by the log2(denominator). The sign bits are the bit-width
// minus this value. The result of this subtraction has to be positive.
unsigned ResBits = TyBits - Denominator->logBase2();
```

Denominator->logBase2() computes the floor of the log, IIUC we should be computing the ceil.





```
1102     ConstantRange
1103     ↘ ConstantRange::add(const ConstantRange &Other) const {
1104         if (isEmptySet() || Other.isEmptySet())
1105             return getEmpty();
1106         if (isFullSet() || Other.isFullSet())
1107             return getFull();
1108
1109         APInt NewLower = getLower() + Other.getLower();
1110         APInt NewUpper = getUpper() + Other.getUpper() - 1;
1111         if (NewLower == NewUpper)
1112             return getFull();
1113
1114         ConstantRange X = ConstantRange(std::move(NewLower), std::move(NewUpper));
1115         if (X.isSizeStrictlySmallerThan(*this) ||
1116             X.isSizeStrictlySmallerThan(Other))
1117             // We've wrapped, therefore, full set.
1118             return getFull();
1119         return X;
1120     }
1121 }
```

[l1, u1], [l2, u2] → [l3, h3]

+ #

```
1102     ConstantRange
1103     ↘ ConstantRange::add(const ConstantRange &Other) const {
1104         if (isEmptySet() || Other.isEmptySet())
1105             return getEmpty();
1106         if (isFullSet() || Other.isFullSet())
1107             return getFull();
1108
1109         APInt NewLower = getLower() + Other.getLower();
1110         APInt NewUpper = getUpper() + Other.getUpper() - 1;
1111         if (NewLower == NewUpper)
1112             return getFull();
1113
1114         ConstantRange X = ConstantRange(std::move(NewLower), std::move(NewUpper));
1115         if (X.isSizeStrictlySmallerThan(*this) ||
1116             X.isSizeStrictlySmallerThan(Other))
1117             // We've wrapped, therefore, full set.
1118             return getFull();
1119         return X;
1120     }
1121 }
```

[l1, u1], [l2, u2] → [l3, h3]

+

#

Overflow

Top

Bot

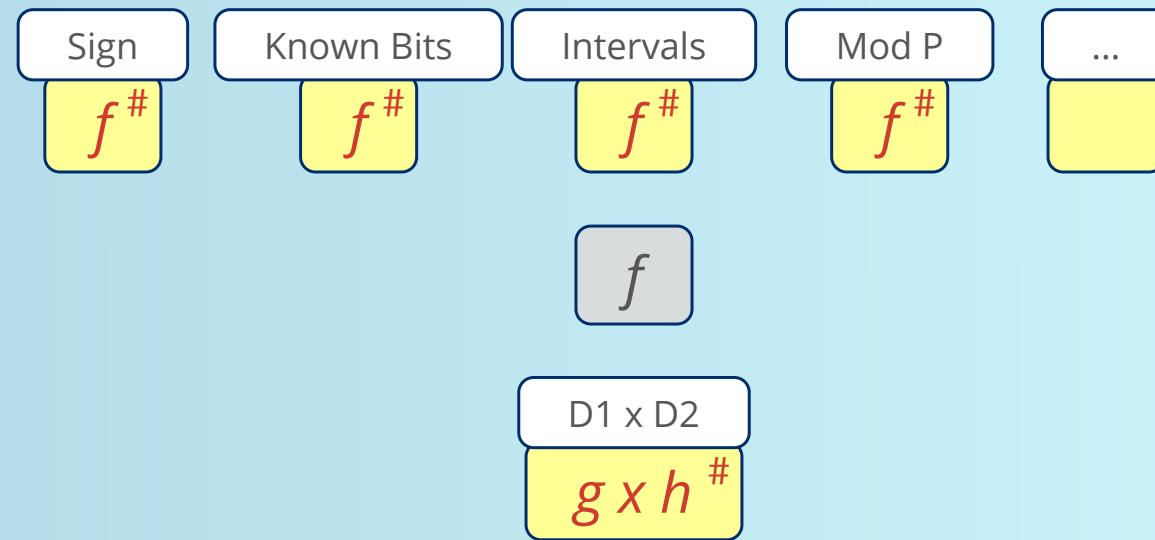
Many Abstract Domains ...

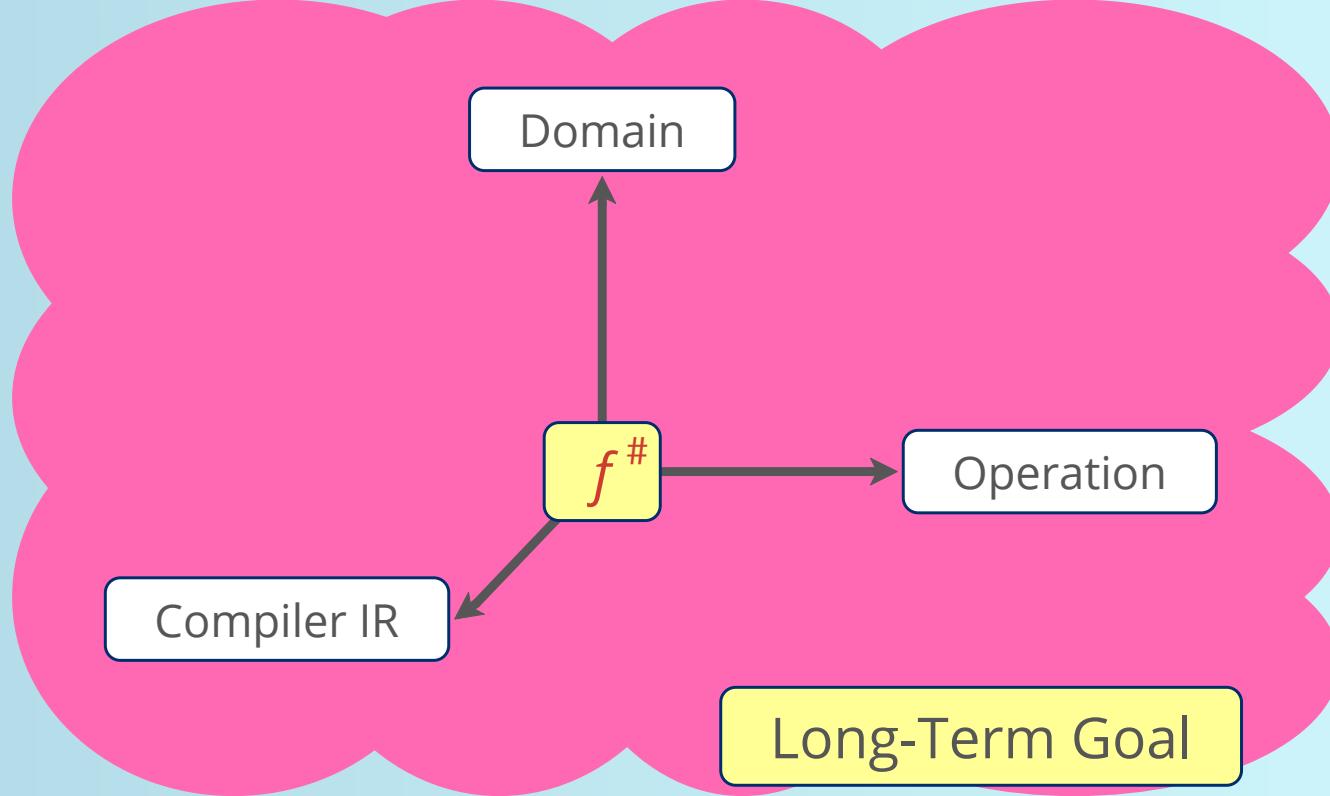
Intervals

$f^\#$

f

Many Abstract Domains ...







Design point: **Bulk Automation**

- Minimize user input
- **Coverage** over precision



Domains = Known Bits, Constant Range

Nice To Meet You

$f^{\#}$

40 binary integer ops

IR = MLIR



Domains = Known Bits, Constant Range

Nice To Meet You

$f^{\#}$

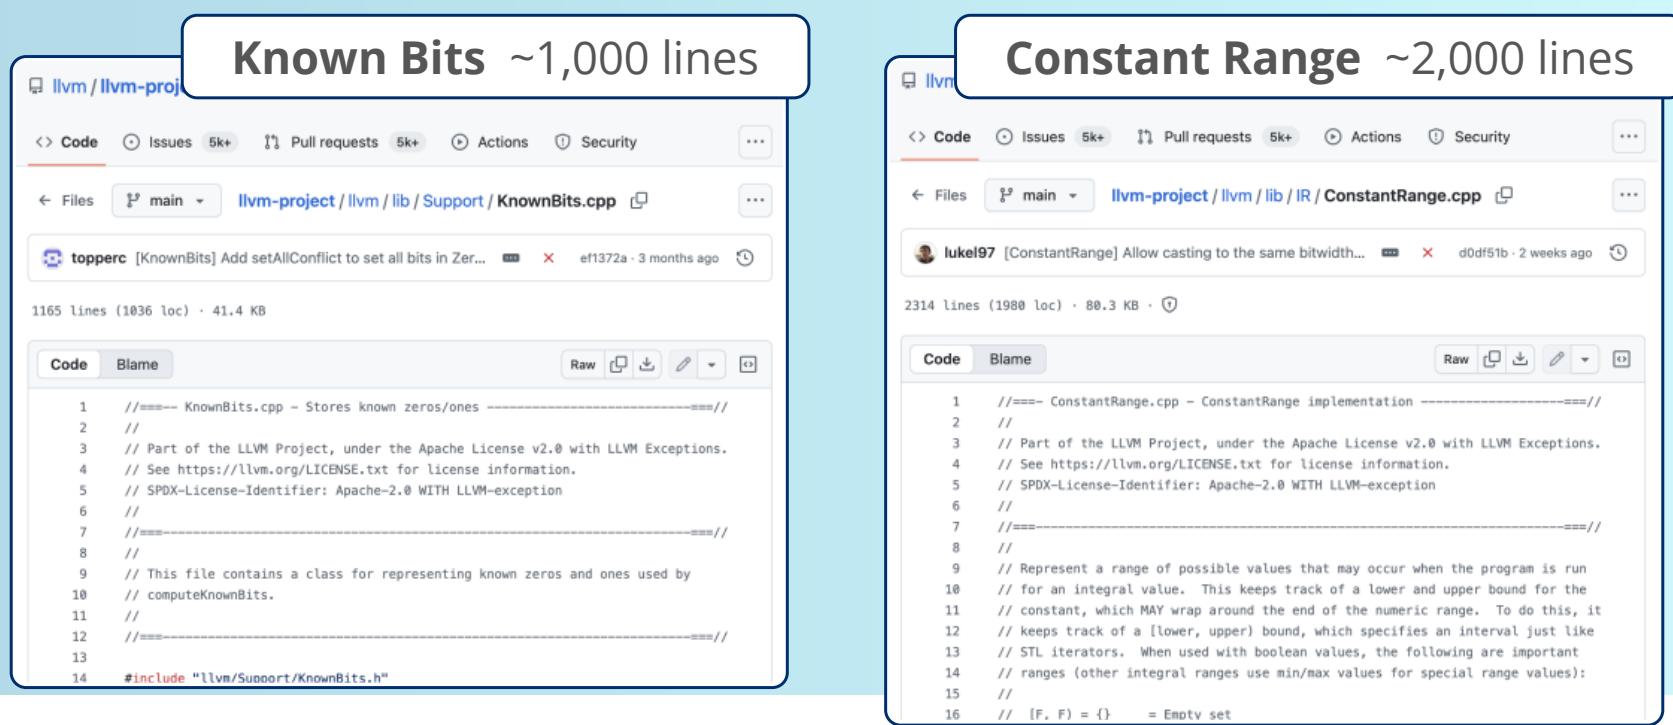
40 binary integer ops

26 ops ~ complementary to LLVM
14 ops ~ more precise

IR = MLIR



Nice To Meet You filled edge-cases in 10+ yr old transformers



Known Bits

Constant Range

Known Bits

Constant Range

Integer intervals $[a, b)$
Signed and Unsigned

$[1, 4)$

1

2

3

Tested on 8-bit and 64-bit integers

Known Bits

Partial bit vectors

10??

8

9

10

11

Constant Range

Integer intervals $[a, b)$
Signed and Unsigned

[1, 4)

1

2

3



... but synthesis doesn't scale?



... but synthesis doesn't scale?



... but synthesis doesn't scale?



Idea : build up a **meet** of several transformers

$$f_1^{\#} \sqcap f_2^{\#} \sqcap \dots \sqcap f_n^{\#}$$

- **dynamically adapt** score function to cover input space

Give each transformer **one simple job**

Add

AddNsw

AddNswNuw

AddNuw

Algorithm 2: MCMCSynthesizeTransformer($\mathcal{F}^s, prob$)

```
1 Input:  $\mathcal{F}^s$  – Current set of synthesized transformers;  
2 Output: A new sound transformer  $f \in \mathcal{L}$  that (in the limit) minimizes precision.  
3 fun Cost( $f^\#$ ):  
4   return  $\lambda(1 - \text{Soundness}(f^\#)) + \kappa(1 - \text{Improvement}(f^\#, \mathcal{F}^s))$  // reward soundness, precision improv.  
5  $f \leftarrow \text{initialize}()$  // random initial program  
6 for  $i \leftarrow 1$  to  $N_{step}$  do  
7    $f' \leftarrow \text{mutate}(f)$  // mutate current candidate  
8    $p \sim \mathcal{U}(0, 1)$  // sample acceptance threshold  
9   if Cost( $f$ ) – Cost( $f'$ )  $> T \cdot \log(p)$  then  
10     $f \leftarrow f'$  // accept proposed candidate  
11 if Soundness( $f$ )  $< 1$  then  
12   return  $\top$  // return trivial top transformer if no sound one found  
13 return  $f$  // return the lowest cost sound transformer found
```

Example: Interval Max

Ideal

$$f^{\#}(a, b) = [\max(a.\text{lo}, b.\text{lo}), \max(a.\text{hi}, b.\text{hi})]$$

Example: Interval Max

Ideal

$$f^{\#}(a, b) = [\max(a.\text{lo}, b.\text{lo}), \max(a.\text{hi}, b.\text{hi})]$$

Synth

$$f1^{\#}(a, b) = [0, \max(a.\text{hi}, b.\text{hi})]$$

$$f2^{\#}(a, b) = [a \& b, \text{MAX_INT}]$$

$$f3^{\#}(a, b) = [a.\text{lo}, a.\text{hi} \mid b.\text{hi}]$$

$$f4^{\#}(a, b) = [b.\text{lo}, \text{MAX_INT}]$$

Example: Interval Max

Input: the Domain and Op

Concretization: $\gamma([a.l, a.r]) = \{a.l, \dots, a.r\}$

Meet: $a \sqcap b = [\max(a.l, b.l), \min(a.r, b.r)]$

Join: $a \sqcup b = [\min(a.l, b.l), \max(a.r, b.r)]$

Abstraction: $\beta(x) = [x, x]$

Concrete op: $f(x, y) = \max(x, y)$

DSL ops: $\{+, -, \&, |, \min, \max, \dots\}$

Size: $|a| = \lfloor \log_2(|a.l - a.r|) \rfloor$

What are we looking for?

$$\mathcal{P}(C) \xrightleftharpoons[\alpha]{\gamma} \mathcal{A}$$

The best abstract transformer:

$$f^{\#} = \alpha \circ f^* \circ \gamma$$

Algorithm 2: MCMCSynthesizeTransformer($\mathcal{F}^s, prob$)

```
1 Input:  $\mathcal{F}^s$  – Current set of synthesized transformers;  
2 Output: A new sound transformer  $f \in \mathcal{L}$  that (in the limit) minimizes precision.  
3 fun Cost( $f^\#$ ):  
4   return  $\lambda(1 - \text{Soundness}(f^\#)) + \kappa(1 - \text{Improvement}(f^\#, \mathcal{F}^s))$  // reward soundness, precision improv.  
5  $f \leftarrow \text{initialize}()$  // random initial program  
6 for  $i \leftarrow 1$  to  $N_{step}$  do  
7    $f' \leftarrow \text{mutate}(f)$  // mutate current candidate  
8    $p \sim \mathcal{U}(0, 1)$  // sample acceptance threshold  
9   if Cost( $f$ ) – Cost( $f'$ )  $> T \cdot \log(p)$  then  
10     $f \leftarrow f'$  // accept proposed candidate  
11 if Soundness( $f$ )  $< 1$  then  
12   return  $\top$  // return trivial top transformer if no sound one found  
13 return  $f$  // return the lowest cost sound transformer found
```

Algorithm 2: MCMCSynthesizeTransformer(\mathcal{F}^s , $prob$)

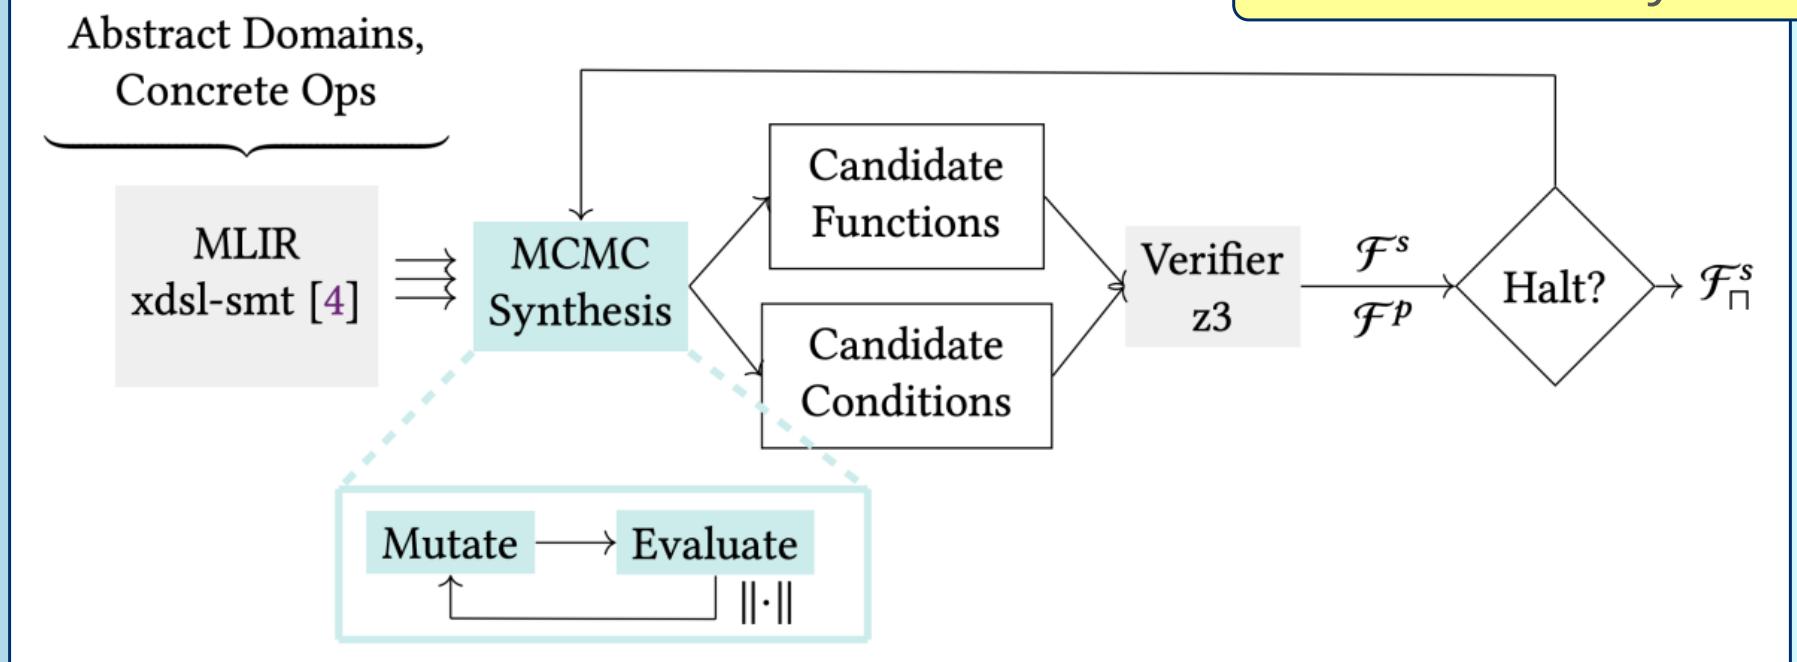
```
1 Input:  $\mathcal{F}^s$  – Current set of synthesized transformers;  
2 Output: A new sound transformer  $f \in \mathcal{L}$  that (in the limit) minimizes precision.  
3 fun Cost( $f^\#$ ):  
4   return  $\lambda(1 - \text{Soundness}(f^\#)) + \kappa(1 - \text{Improvement}(f^\#, \mathcal{F}^s))$  // reward soundness, precision improv.  
5  $f \leftarrow \text{initialize}()$  // random initial program  
6 for  $i \leftarrow 1$  to  $N_{step}$  do  
7    $f' \leftarrow \text{mutate}(f)$  // mutate current candidate  
8    $p \sim \mathcal{U}(0, 1)$  // sample acceptance threshold  
9   if Cost( $f$ ) – Cost( $f'$ )  $> T \cdot \log(\frac{p}{1-p})$  then  
10     $f \leftarrow f'$  // accept proposed candidate  
11 if Soundness( $f$ )  $< 1$  then  
12   return  $\top$  // return trivial top transformer if no sound one found  
13 return  $f$  // return the lowest cost sound transformer found
```

Optimal in the limit
mutate is invertible

Sound at each step

Nice To Meet You

1. Eval + Verify
2. Condition Synth



<https://github.com/dominicmkennedy/synth-xfer>



MLIR

<https://mlir.llvm.org/>

Multi-Level IR Compiler Framework

SSA : Single Static Assignment

One syntax, many **dialects**



MLIR

<https://mlir.llvm.org/>

Multi-Level IR Compiler Framework

SSA : Single Static Assignment

One syntax, many **dialects**

Our dialect: **xDSL-SMT**



U. Cambridge

<https://github.com/opencompl/xdsl-smt>

CGO'25

PLDI'25

Example: MLIR vs. SMT-LIB

$x * 2 \neq x + x$

```
%x = smt.declare_const : !smt_int.int
%two = smt_int.constant 2
%mul_two = smt_int.mul %x, %two
%add_twice = smt_int.add %x, %x
%eq = smt.distinct %mul_two, %add_twice
      : !smt_int.int
smt.assert %eq
smt.check_sat
```

```
(declare-const x Int)
(assert
  (distinct (* x 2) (+ x x)))
(check-sat)
```

Synthesized urem

```
func.func @f1(%L : KnownBits, %R : KnownBits) -> KnownBits {
    %1 = countLeadingZero(%L.zero)
    %knownZero = setHighBits(0, %1)
    return makeKnownBits(%knownZero, 0)
}
func.func @f2_cond(%L: KnownBits, %R: KnownBits) -> bool {
    %Lmax = negate(%L.zero)
    %Rmin = %R.one
    %cond = unsignedLessThan(%Lmax, %Rmin)
    return %cond
}
func.func @f2_body(%L : KnownBits, %R : KnownBits) -> KnownBits {
    return %L
}
func.func @f2(%L : KnownBits, %R : KnownBits) -> KnownBits {
    return ite(@f2_cond(%L, %R), @f2_body(%L, %R), %top)
}
...
func.func @solution(%L : KnownBits, %R : KnownBits) -> KnownBits {
    return meet(@f1(%L, %R), ... , @f9(%L, %R))
}
```

Synthesized urem

```
func.func @f1(%L : KnownBits, %R : KnownBits) -> KnownBits {  
    %1 = countLeadingZero(%L.zero)  
    %knownZero = setHighBits(0, %1)  
    return makeKnownBits(%knownZero, 0)  
}  
func.func @f2_cond(%L: KnownBits, %R: KnownBits) -> bool {  
    %Lmax = negate(%L.zero)  
    %Rmin = %R.one  
    %cond = unsignedLessThan(%Lmax, %Rmin)  
    return %cond  
}  
func.func @f2_body(%L : KnownBits, %R : KnownBits) -> KnownBits {  
    return %L  
}  
func.func @f2(%L : KnownBits, %R : KnownBits) -> KnownBits {  
    return ite(@f2_cond(%L, %R), @f2_body(%L, %R), %top)  
}  
...  
func.func @solution(%L : KnownBits, %R : KnownBits) -> KnownBits {  
    return meet(@f1(%L, %R), ... , @f9(%L, %R))  
}
```

Meet of 9 candidates

Recovers LLVM heuristics

f2 improves precision

Evaluation 1: Per-Operator Precision

Complementary on 26 / 40 ops

Known Bits

ConcreteOp	#f [#]			Tests				64-bit precision (norm) ↓					
	#f [#]	#c	#inst	T	synth	llvm	meet	T	synth	llvm	meet		
Abds	10	3	189	1000	33.90	60.10	100.00	100.00	10000	0.059	0.050	0.000	0.000
Abdu	16	4	259	1000	33.10	59.40	100.00	100.00	10000	0.059	0.050	0.000	0.000
Add	17	2	306	1000	29.60	58.70	100.00	100.00	10000	0.140	0.082	0.000	0.000
AddNsw	13	2	205	1000	24.50	42.00	100.00	100.00	9674	0.147	0.136	0.000	0.000
AddNswNuw	14	3	220	1000	7.40	45.50	100.00	100.00	7479	0.160	0.136	0.000	0.000
AddNuw	17	4	291	1000	15.20	53.90	100.00	100.00	8305	0.152	0.103	0.000	0.000

Constant Range

ConcreteOp	#f [#]			Tests				64-bit precision (norm) ↓					
	#f [#]	#c	#inst	T	synth	llvm	meet	T	synth	llvm	meet		
Abds*	3	2	70	1000	59.80	59.80	N/A	59.80	10000	0.917	0.915	N/A	0.915
Abdu	20	6	344	1000	0.00	75.00	N/A	75.00	10000	0.990	0.908	N/A	0.908
Add	2	2	45	509	36.54	100.00	100.00	100.00	4991	0.949	0.887	0.887	0.887
AddNsw*	8	4	148	1000	7.10	100.00	100.00	100.00	9770	0.982	0.905	0.905	0.905
AddNswNuw	10	2	172	1000	0.00	70.70	84.80	88.60	8190	0.994	0.921	0.912	0.910
AddNuw	14	5	243	1000	0.00	94.00	100.00	100.00	8267	0.993	0.910	0.906	0.906

Evaluation 2: SPEC CPU 2017 Known Bits

Evaluation 2: SPEC CPU 2017 Known Bits

Meet of synth & LLVM transformers

	# KB	Baseline KB
openssl	+2	1.3 M
ffmpeg	+14	3.7 M
cvc5	+0	16 M

Evaluation 2: SPEC CPU 2017 Known Bits

Synth **vs.** LLVM-- *Synth always loses!*

	Precision	Time
perlbench	-3.76%	-1.5s
gcc	-1.78%	-6s
mcf	0	-0.2s
omnetpp	-0.24%	-0.1s
xalancbmk	-6.92%	-1s
x264	-11.79%	-6s
deepsjeng	-5.12%	-0.3s
leela	-23.22%	-0.2s
xz	-6.26%	-0.5s

Evaluation 3: Product: Known Bits x Constant Range

Concrete Op	Tests	8-bit exact (%) ↑			Tests	64-bit precision (norm) ↓		
		T	synth	reduced		T	synth	reduced
Abds	1000	7.64	18.06	28.98	10000	0.1260	0.1119	0.1069
Abdu	1000	7.11	20.76	71.06	10000	0.1236	0.1085	0.0921
AddNsw	1000	6.68	18.73	81.65	10000	0.2820	0.1125	0.0869
AddNswNuw	1000	0.22	44.89	88.20	10000	0.5592	0.1216	0.0610
AddNuw	1000	3.35	31.80	92.32	10000	0.4920	0.0930	0.0557
AvgCeilS	1000	9.83	18.01	29.16	10000	0.1651	0.1573	0.1503
AvgFloorS	1000	9.86	17.37	46.61	10000	0.1669	0.1598	0.1412
AvgFloorU	1000	9.90	19.00	50.37	10000	0.1668	0.1481	0.1336
Sdiv	1000	17.99	27.85	45.61	10000	0.7262	0.2493	0.2229
Smax	1000	0.44	59.89	83.19	10000	0.4959	0.0926	0.0822
Smin	1000	0.43	59.62	84.23	10000	0.4954	0.0953	0.0843
Srem	1000	13.14	22.41	26.92	10000	0.1845	0.1701	0.1689
SshlSat	1000	4.08	33.43	43.35	10000	0.9542	0.3211	0.3148
SubNswNuw	1000	0.33	39.01	77.20	10000	0.5617	0.1299	0.0701
SubNuw	1000	3.52	36.84	90.94	10000	0.4822	0.0763	0.0466
UaddSat	1000	4.11	61.03	83.09	10000	0.4482	0.0874	0.0469
Udiv	1000	0.00	68.66	75.28	10000	0.9845	0.0134	0.0067
UdivExact	1000	0.02	3.21	5.78	10000	1.0000	0.0272	0.0195
Umax	1000	0.54	95.28	99.74	10000	0.4947	0.0016	0.0001
Umin	1000	0.56	92.99	99.59	10000	0.4964	0.0023	0.0003
Urem	1000	2.12	61.45	66.53	10000	0.2677	0.0393	0.0367
UsubSat	1000	4.06	56.03	73.09	10000	0.4508	0.1106	0.0700

Evaluation 4: Specialization

Observation: transformers compose poorly

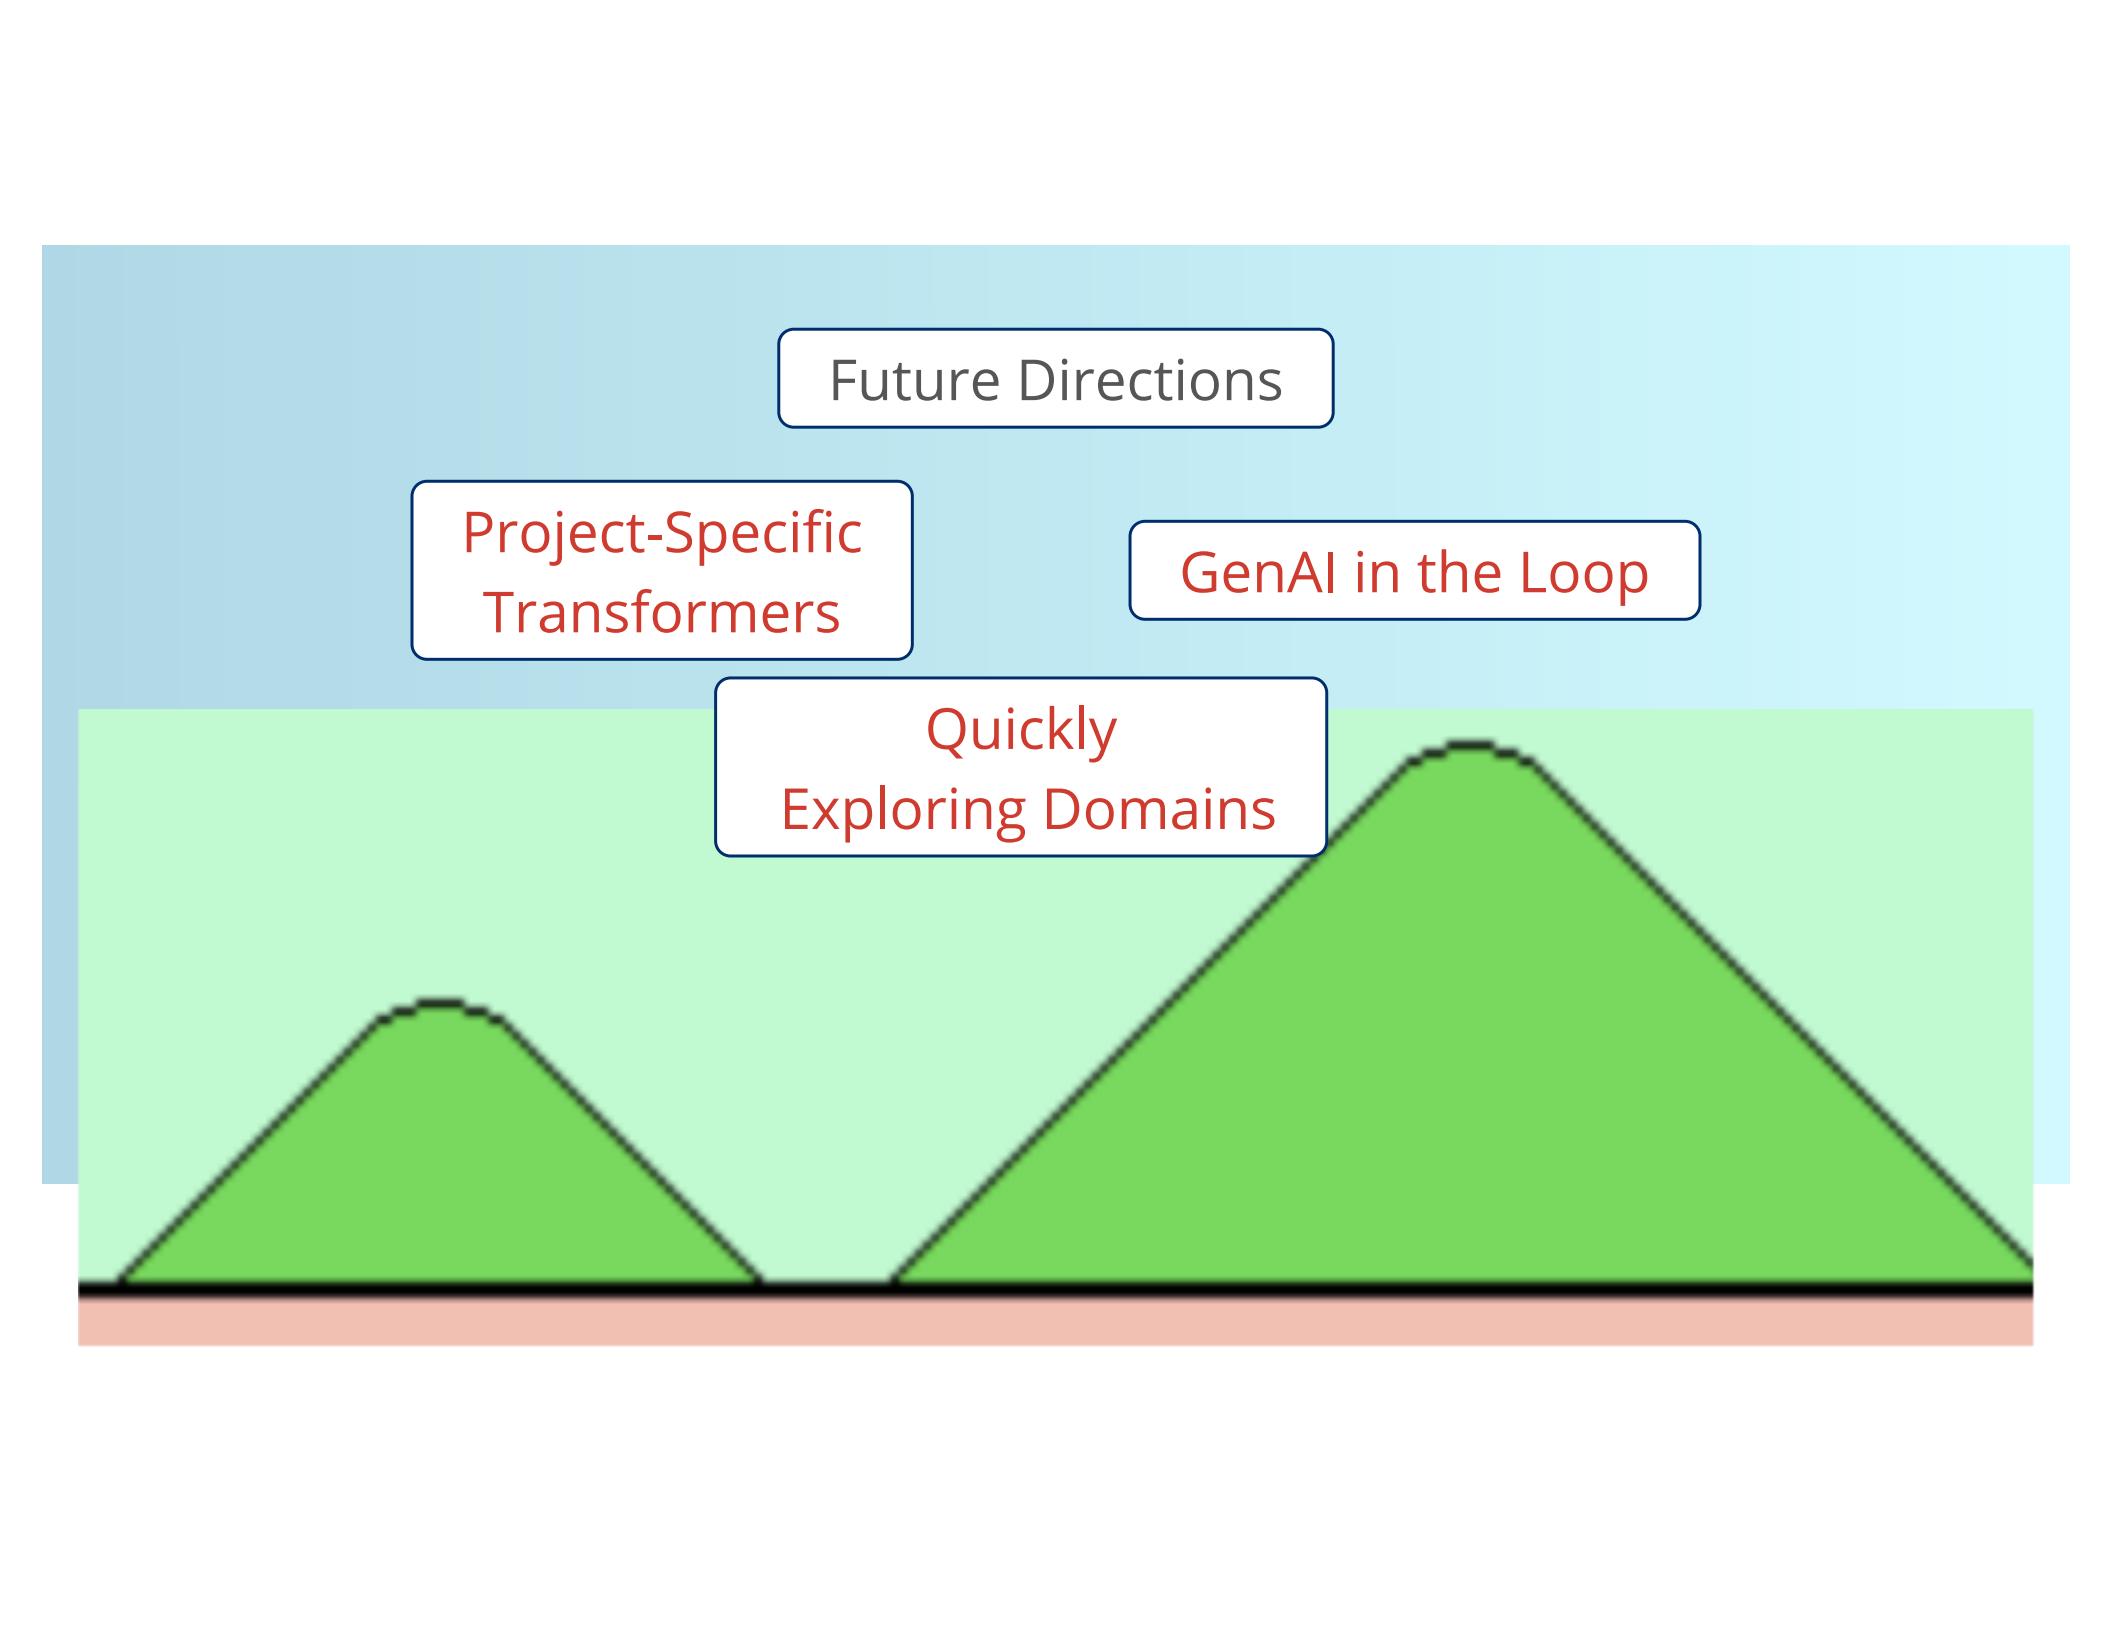
$$f_2^{\#} \circ f_1^{\#} < (f_1 + f_2)^{\#}$$



Evaluation 4: Specialization



Category	Concrete Op	Exact (%) ↑			Precision (norm) ↓		
		T	composed	synth	T	composed	synth
<i>Unary Functions</i> (6,561 test cases)	Abs	1.95	3.92	100.00	3372	2552	0
	CountRZero	0.00	33.33	83.63	5740	3553	193
	CountLZero	0.00	0.00	83.63	5740	5466	193
	PopCount	0.00	0.05	69.53	4461	4456	369
<i>Binary Functions</i> (43,046,721 test cases)	Smax	4.46	6.33	56.86	19,797,600	18,179,000	5,471,520
	Smin	4.46	6.04	70.39	19,797,600	18,384,100	4,117,500
	UaddSat	13.93	22.93	56.20	17,358,900	14,111,200	4,471,530
	UsubSat	13.93	19.46	49.11	17,358,900	15,377,400	5,369,170



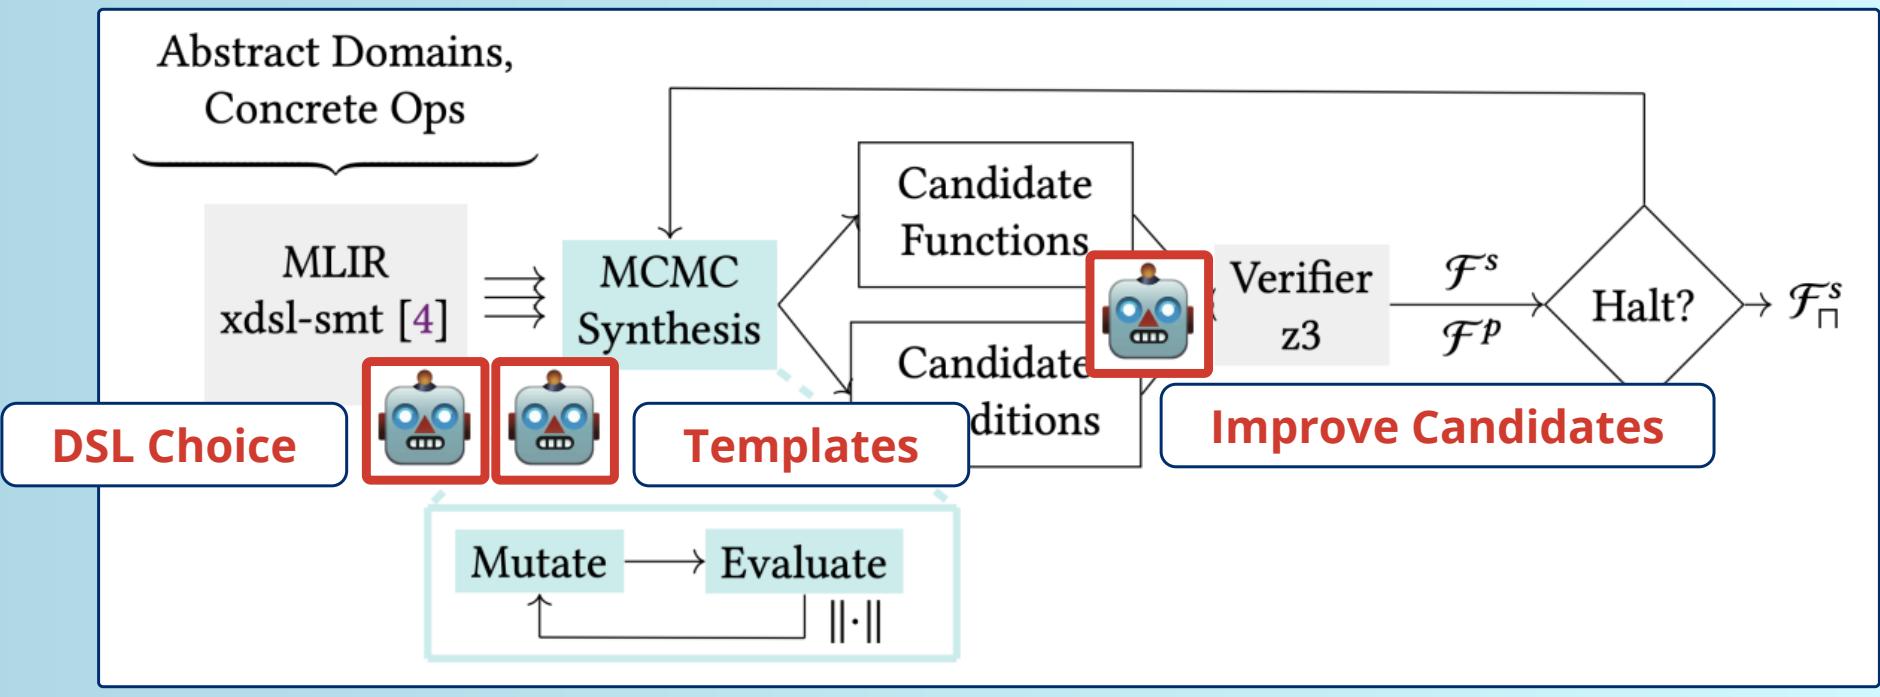
Future Directions

Project-Specific
Transformers

GenAI in the Loop

Quickly
Exploring Domains

GenAI in the Loop



Domains = Known Bits, Constant Range

Nice To Meet You

$f^\#$

40 binary integer ops

IR = MLIR



Meet of sound
candidates

Dynamic cost function

Condition abduction

MLIR + z3 Verification



Related Work

Synthesizing Abstract Transformers **OOPSLA'22**

Kalita, Muduli, D'Antoni, Reps, Roy

Automatic Synthesis of Abstract Operators for eBPF **eBFP'25**

Vishwanathan, Shachnai, Narayana, Nagarakatte

Synthesizing Sound and Precise Abstract Transformers

for Nonlinear Hyperbolic PDE Solvers **OOPSLA'25**

Laurel, Laguna, Huckelheim

Cost-Driven Synthesis of Sound Abstract Interpreters **arxiv'25**

Gu, Singh, Singh

Definition 2.3 (Transformer Synthesis Problem). *Given a concrete transformer $f : C^k \rightarrow C$, an abstract domain $(\mathcal{A}, \top, \gamma, \sqcap, \sqcup, \beta)$, a norm function $\|\cdot\| : (\mathcal{A}^k \rightarrow \mathcal{A}) \rightarrow \mathbb{N}$, and a DSL \mathcal{L} , the transformer synthesis problem is to find a set of transformers $\mathcal{F} = \{f_1^\#, f_2^\#, \dots, f_n^\#\}$ in \mathcal{L} such that*

- Their meet \mathcal{F}_\sqcap is sound: $\text{sound}(\mathcal{F}_\sqcap)$.
- The norm of \mathcal{F}_\sqcap is minimal, i.e., there is no sound set of transformers \mathcal{G} such that $\|\mathcal{G}_\sqcap\| < \|\mathcal{F}_\sqcap\|$.
- No $f_i^\# \in \mathcal{F}$ is redundant: $\forall f_i^\# \in \mathcal{F}, \exists \vec{a} \in \mathcal{A}^k, \left(\prod_{f^\# \in \mathcal{F} \setminus \{f_i^\#\}} f^\#(\vec{a}) \right) \not\subseteq f_i^\#(\vec{a})$.

